Optimizing DTW-Based Audio-to-MIDI Alignment and Matching

Colin Raffel and Daniel P. W. Ellis

41st IEEE International Conference on Acoustics, Speech and Signal Processing
March 23, 2016

Dynamic Time Warping

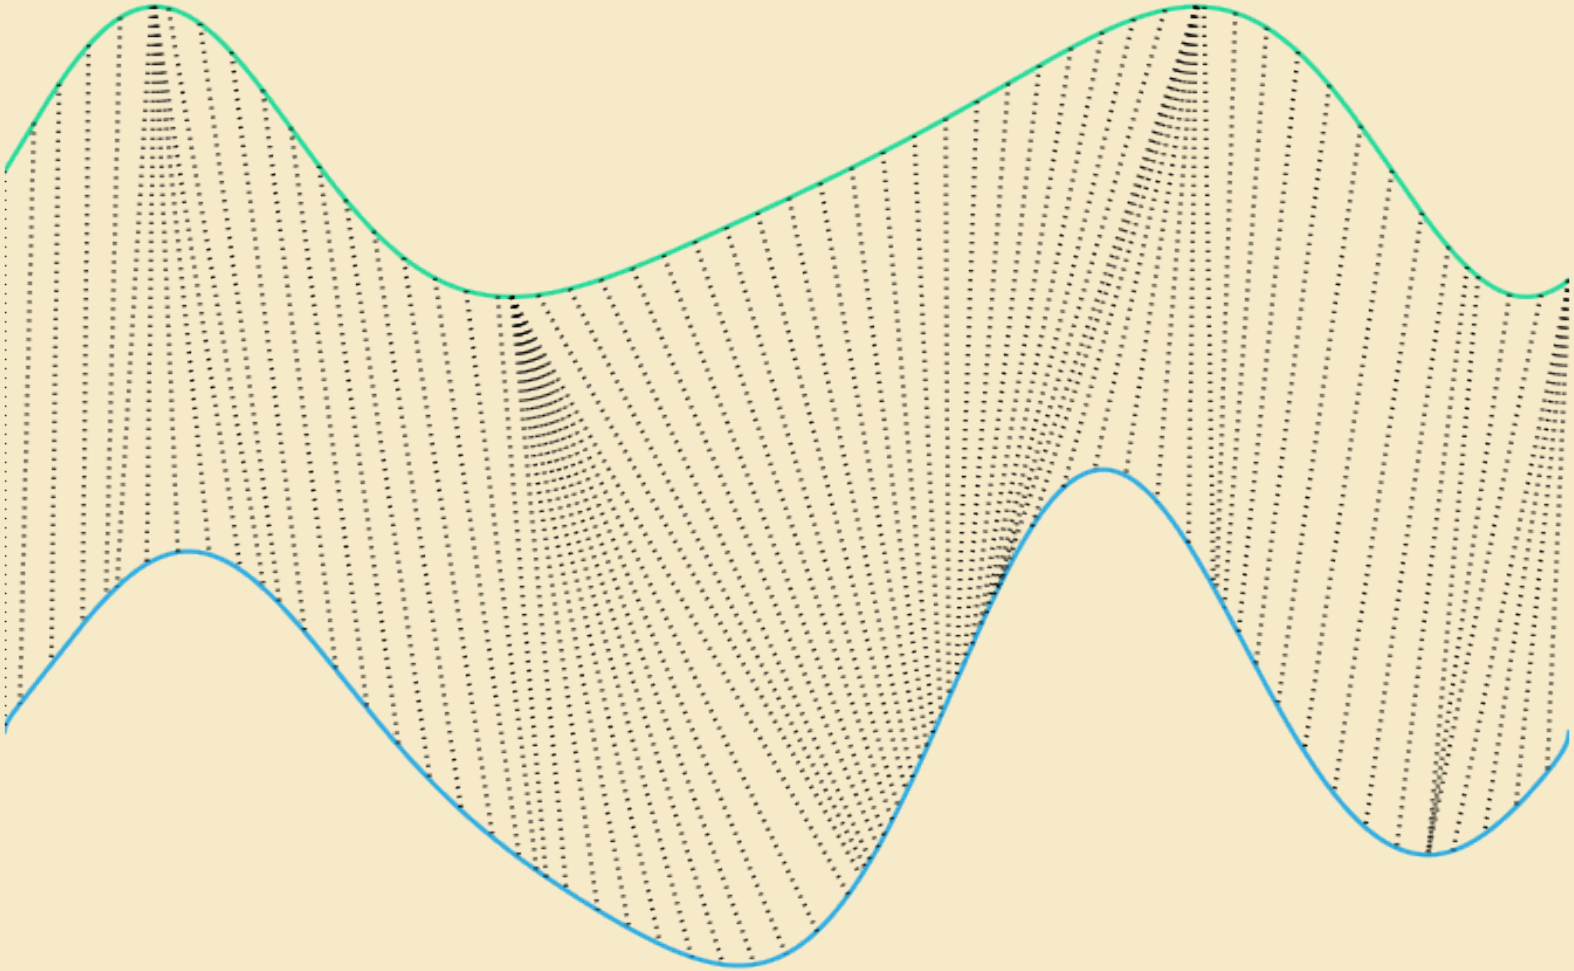

DTW for Audio-to-MIDI Alignment

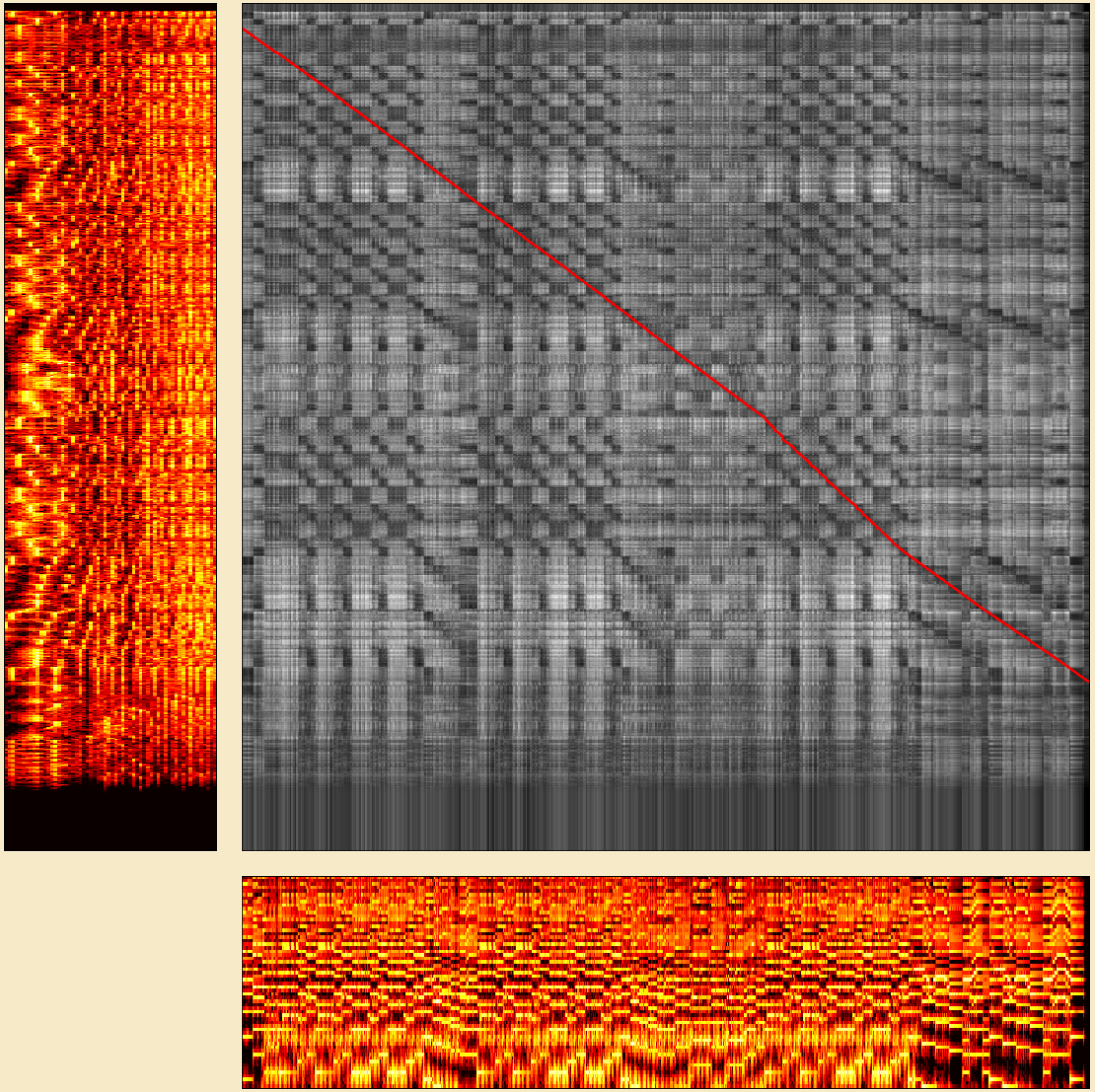

System Design: Representation?

Chromagram?

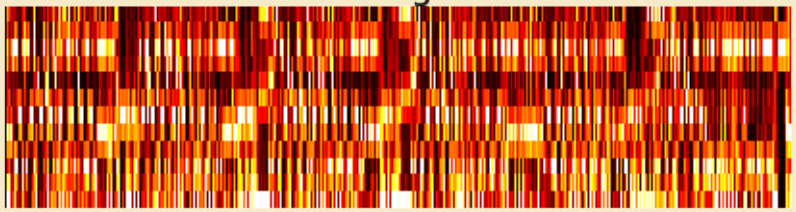

Constant-Q Spectrogram?

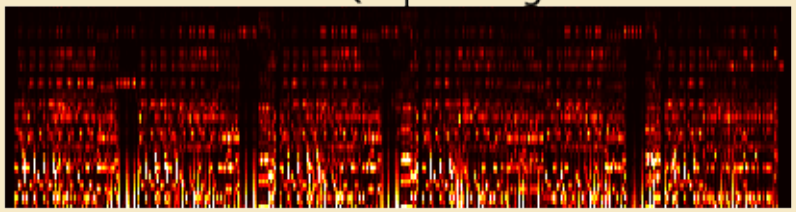

Log Magnitude?

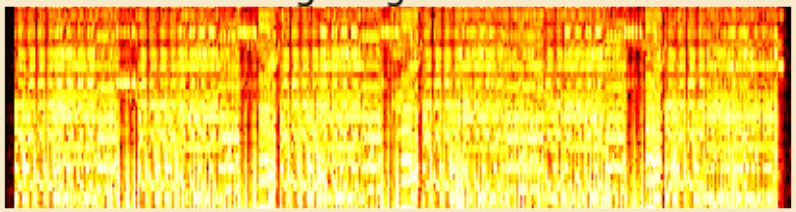

Z-scored?

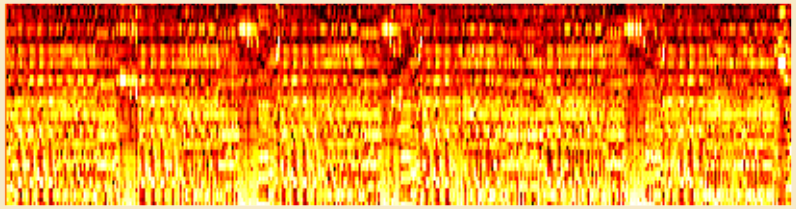

L2-normalized?

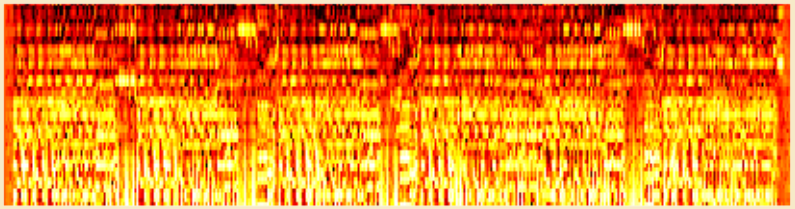

Beat-synchronous?

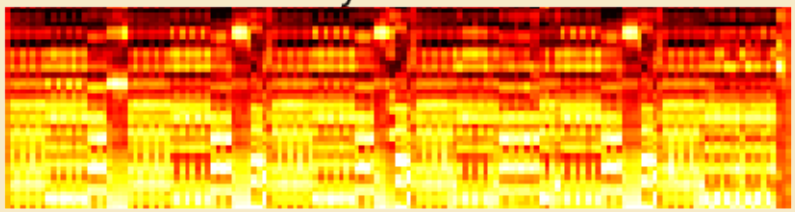

System Design: Path Constraints?

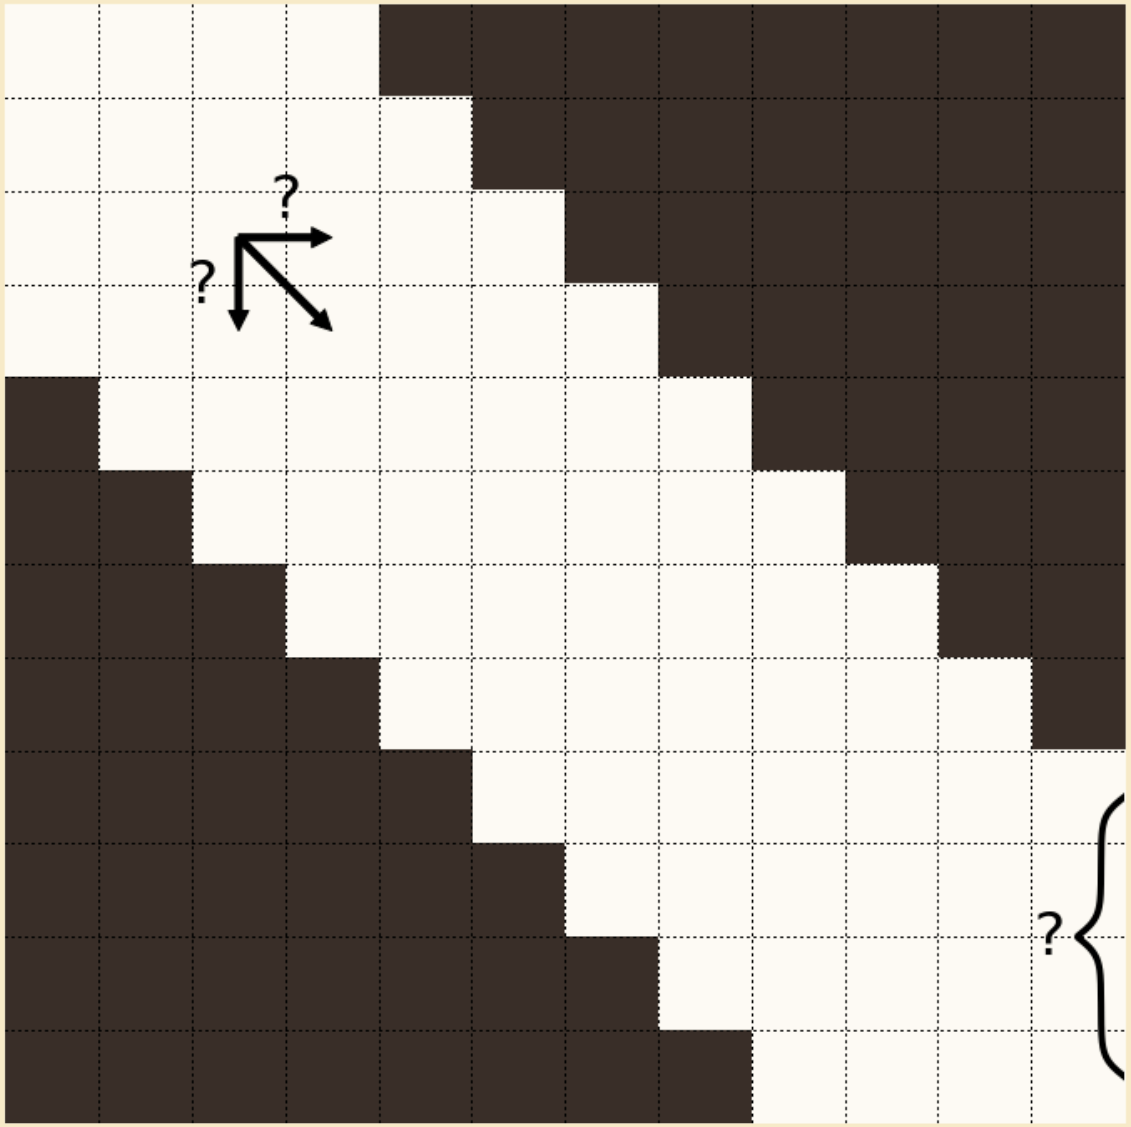

System Design: Score Reporting?

$$\text{score} = \frac{\sum_{i=1}^{|p_m|} D[p_m[i], p_a[i]] + \Phi(i)}{|p_m| \frac{\sum_{i=\min(p_m)}^{\max(p_m)} \sum_{j=\min(p_a)}^{\max(p_a)} D[i,j]}{|\max(p_m) - \min(p_m)||\max(p_a) - \min(p_a)|}}$$

Bayesian Optimization

Idea: Synthetic Alignment Data

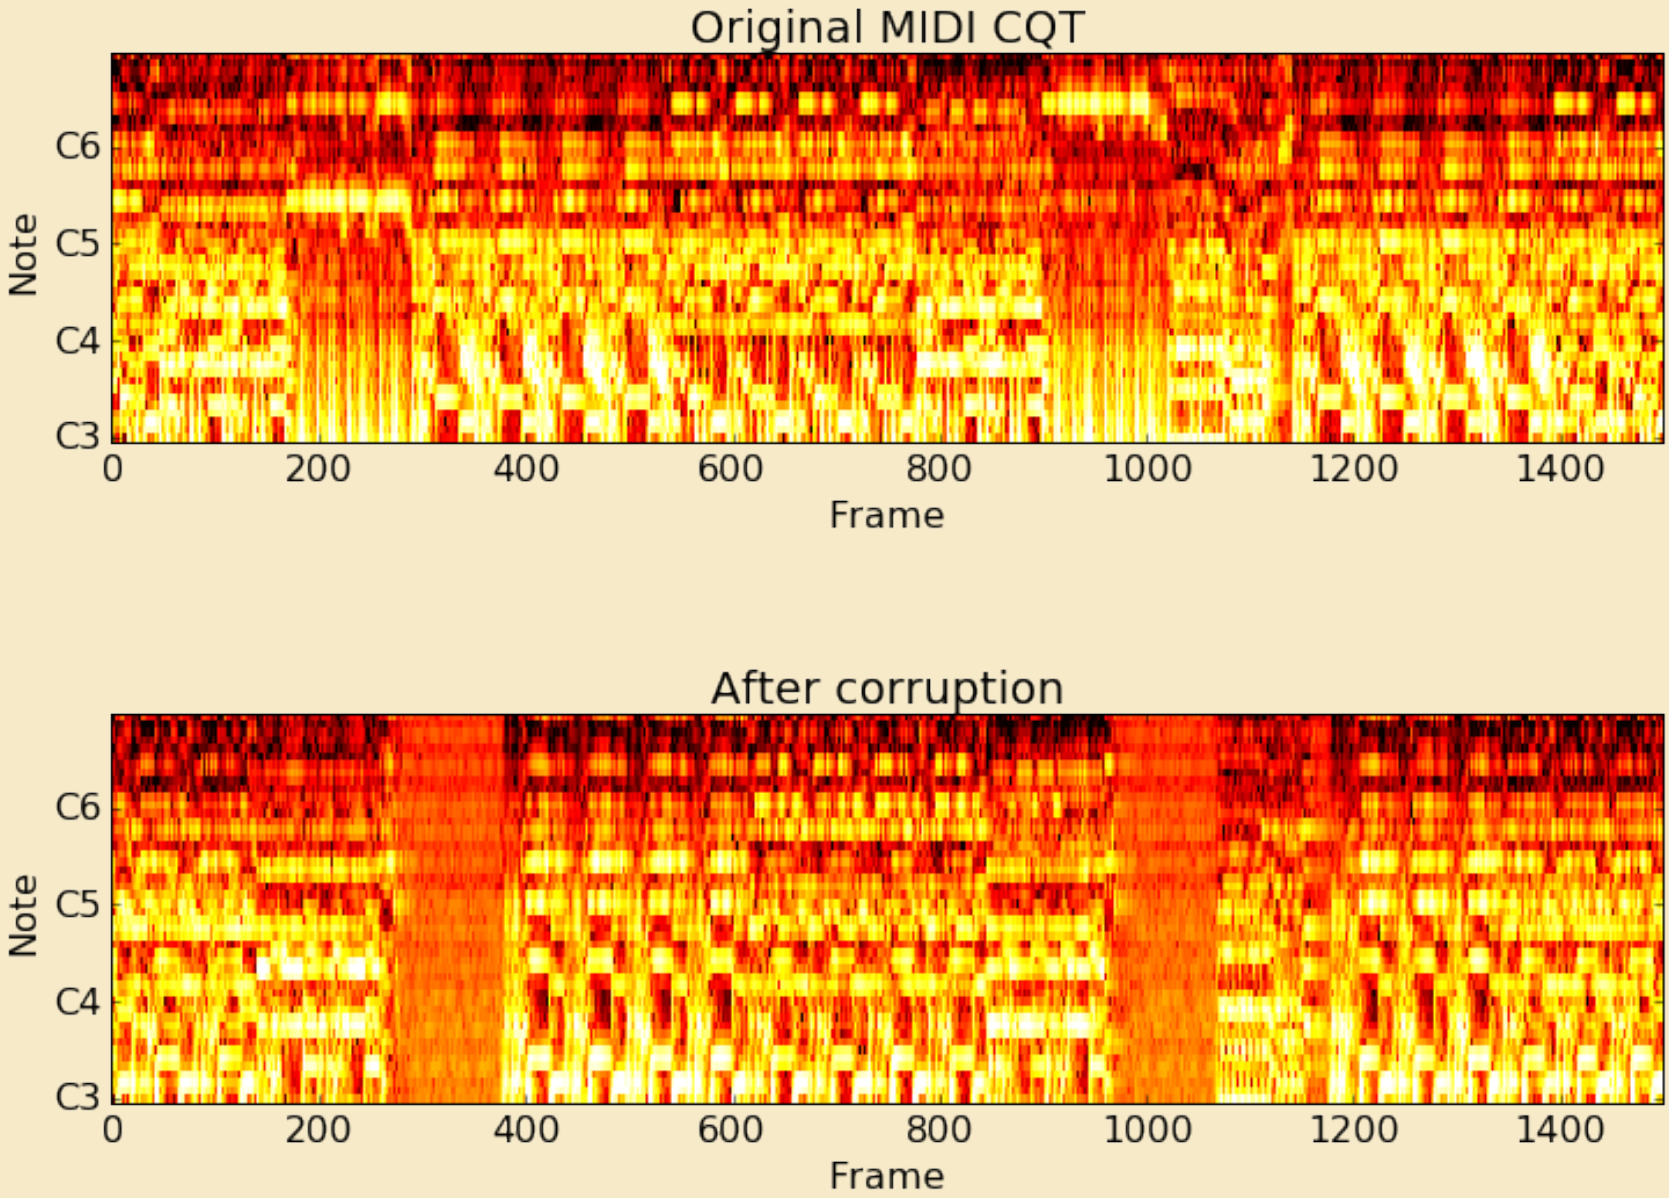

Artificial Time Warping

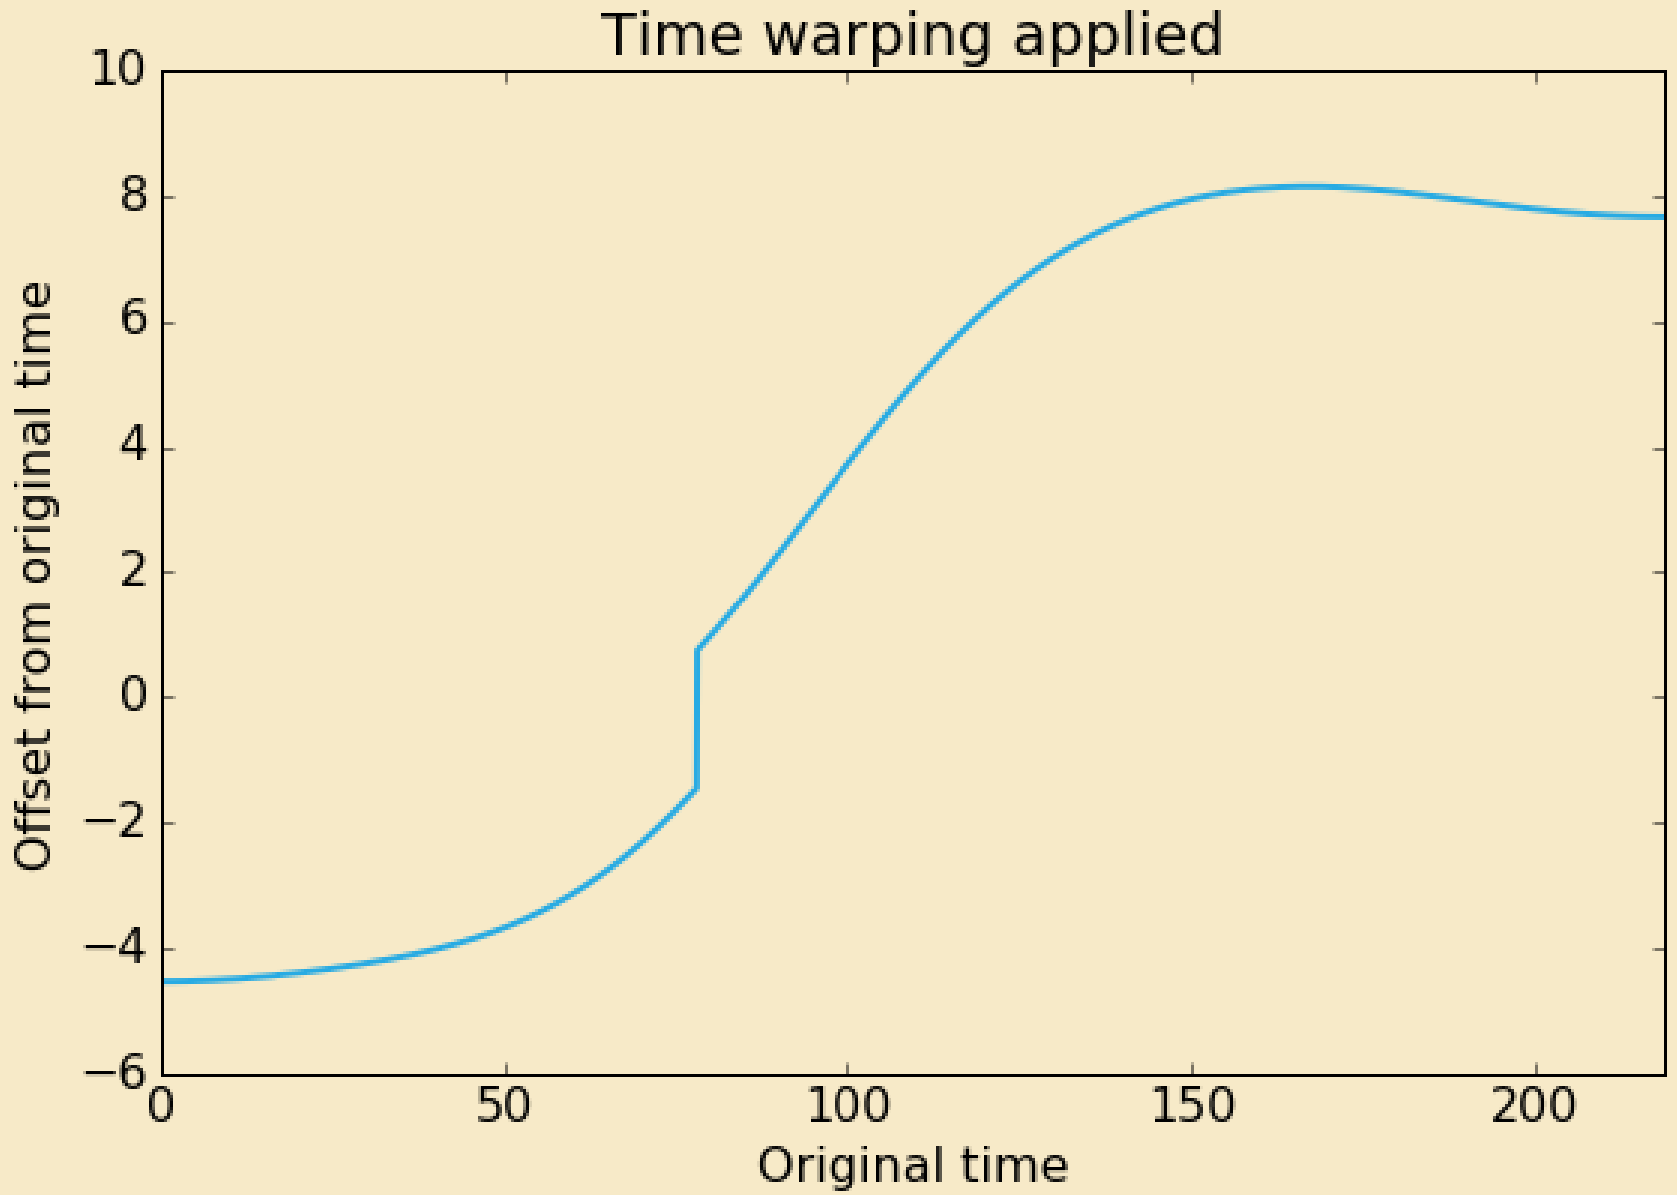

Correcting Time Warping

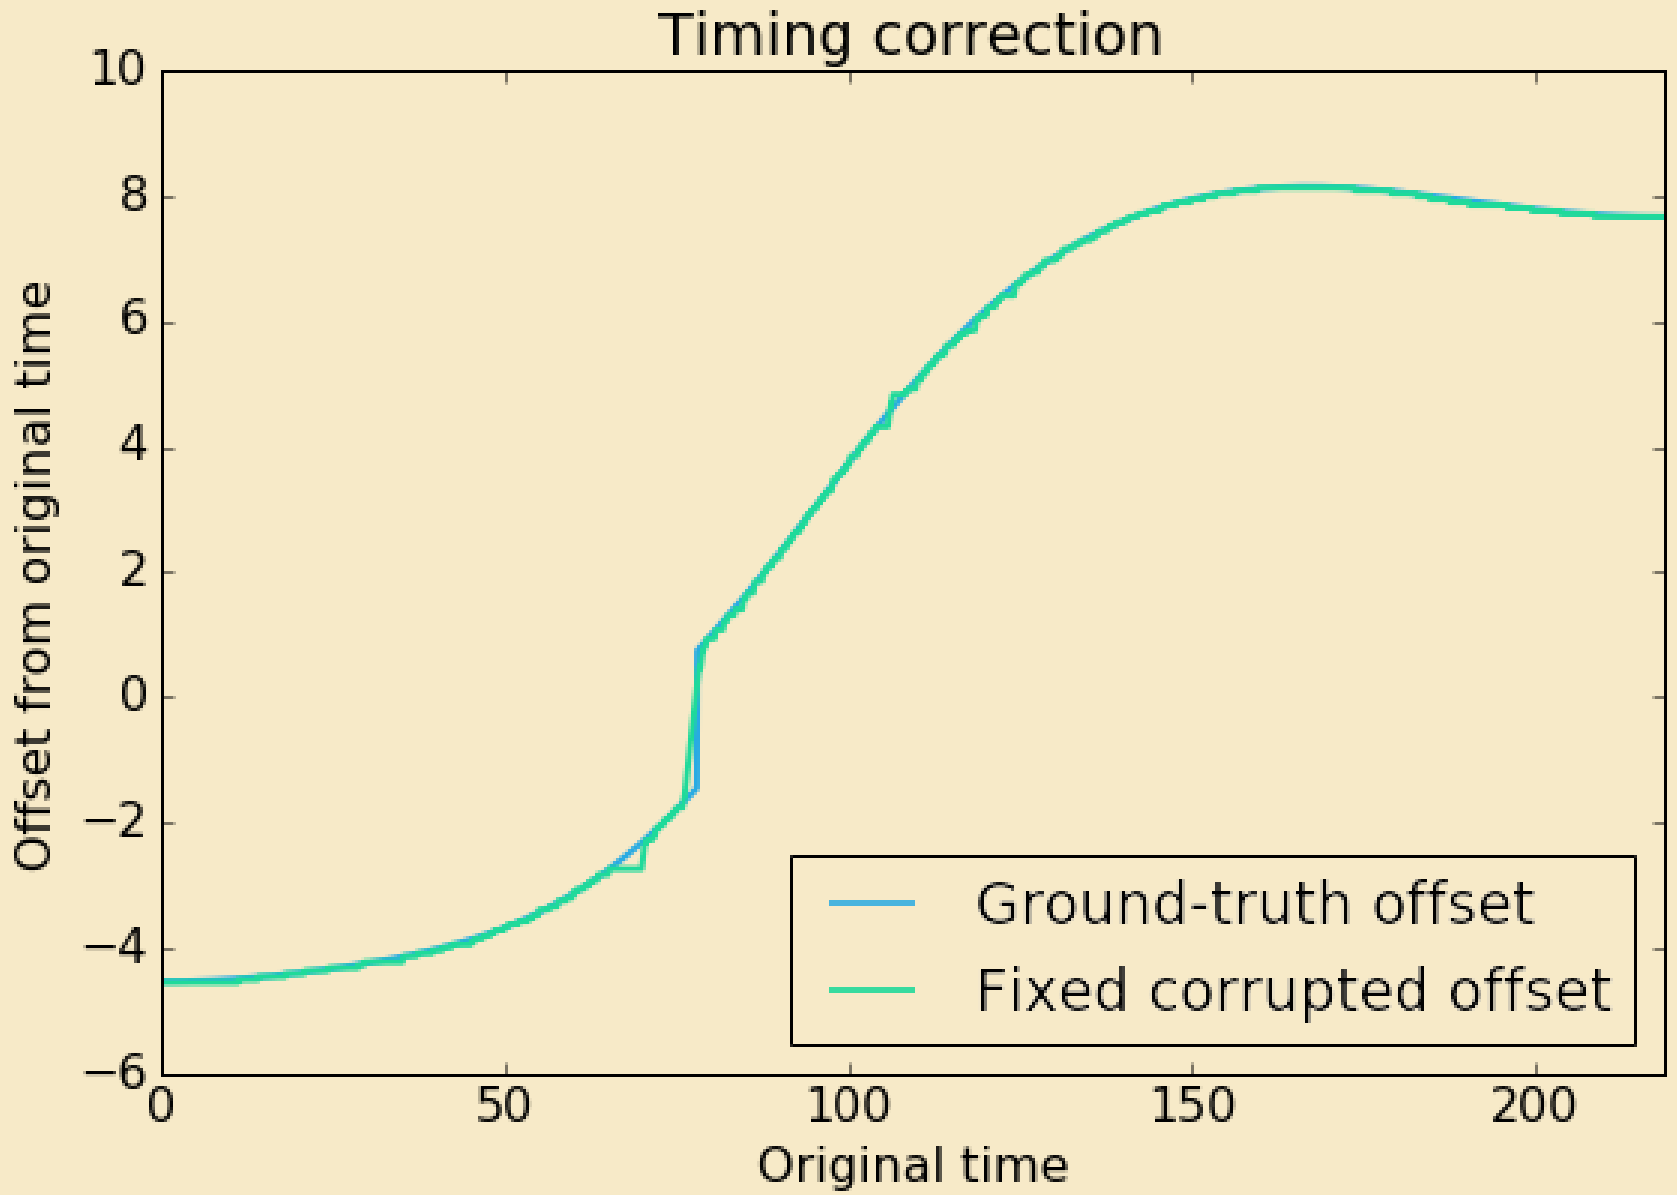

Measuring Error

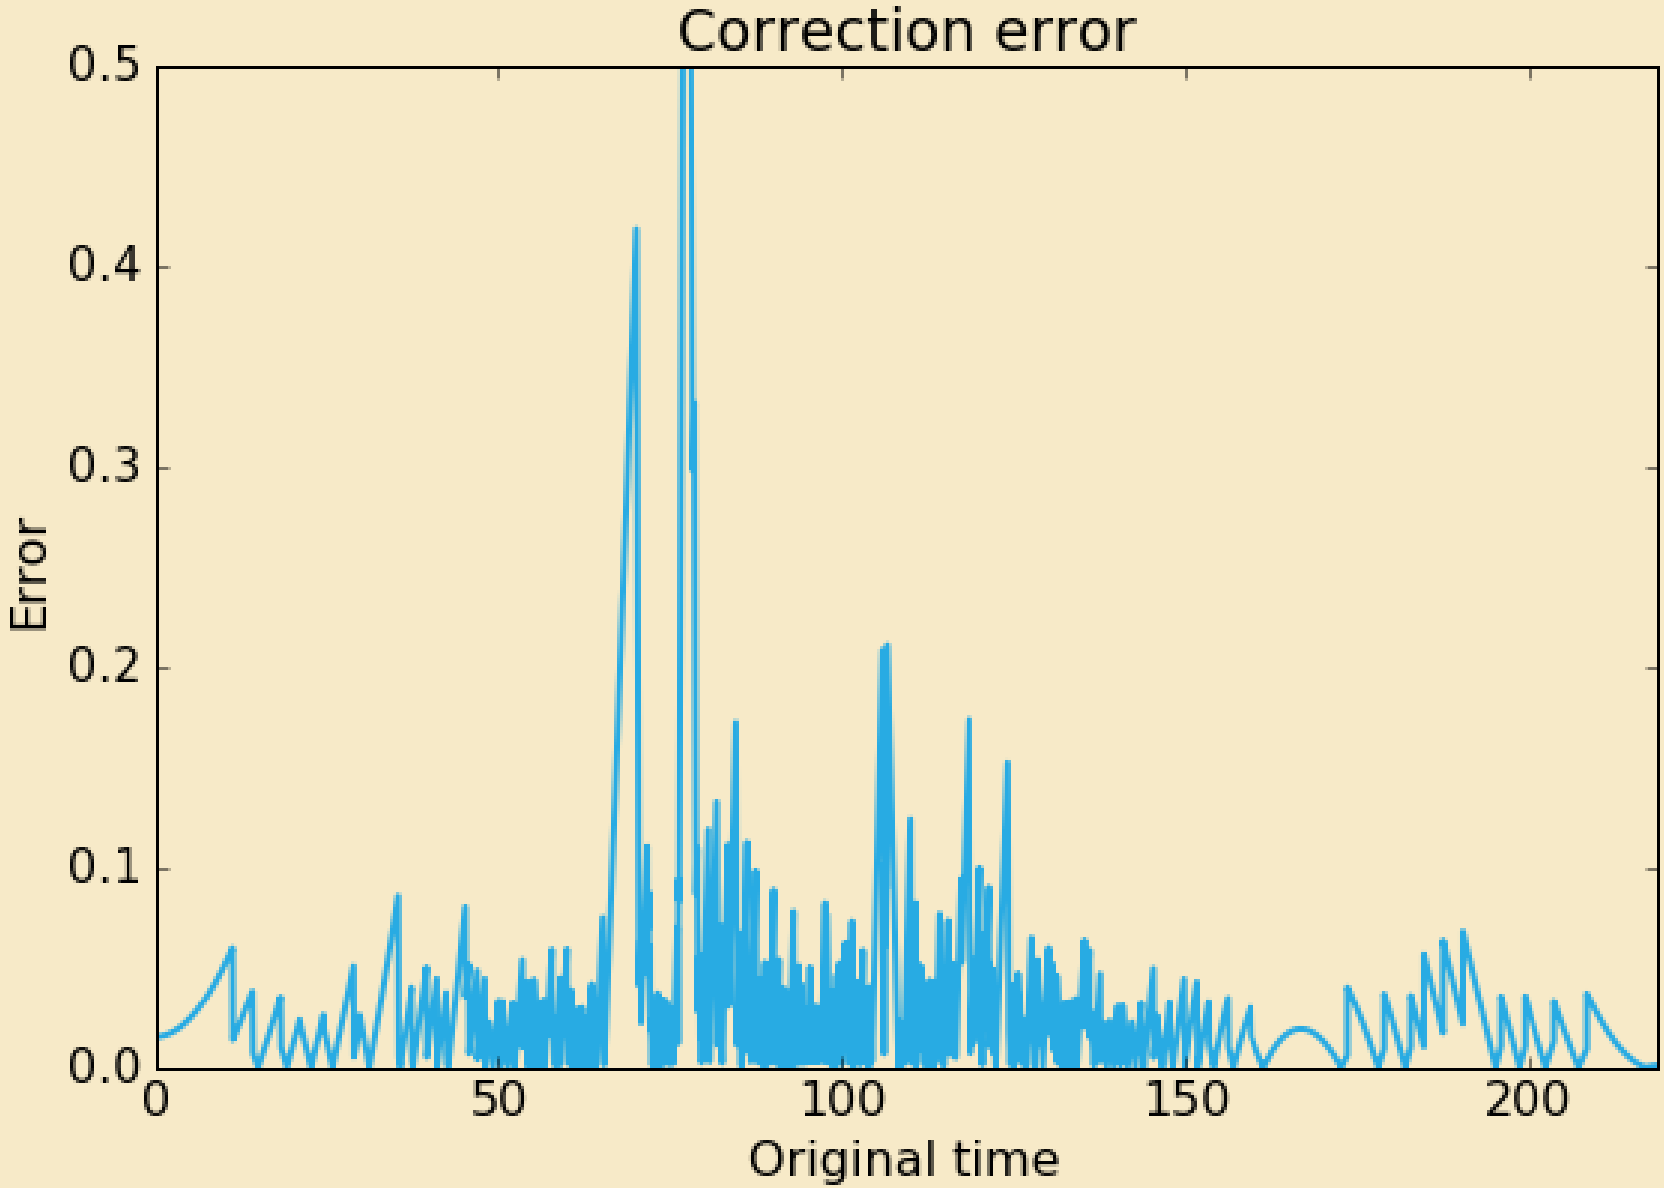

Score Normalization Search

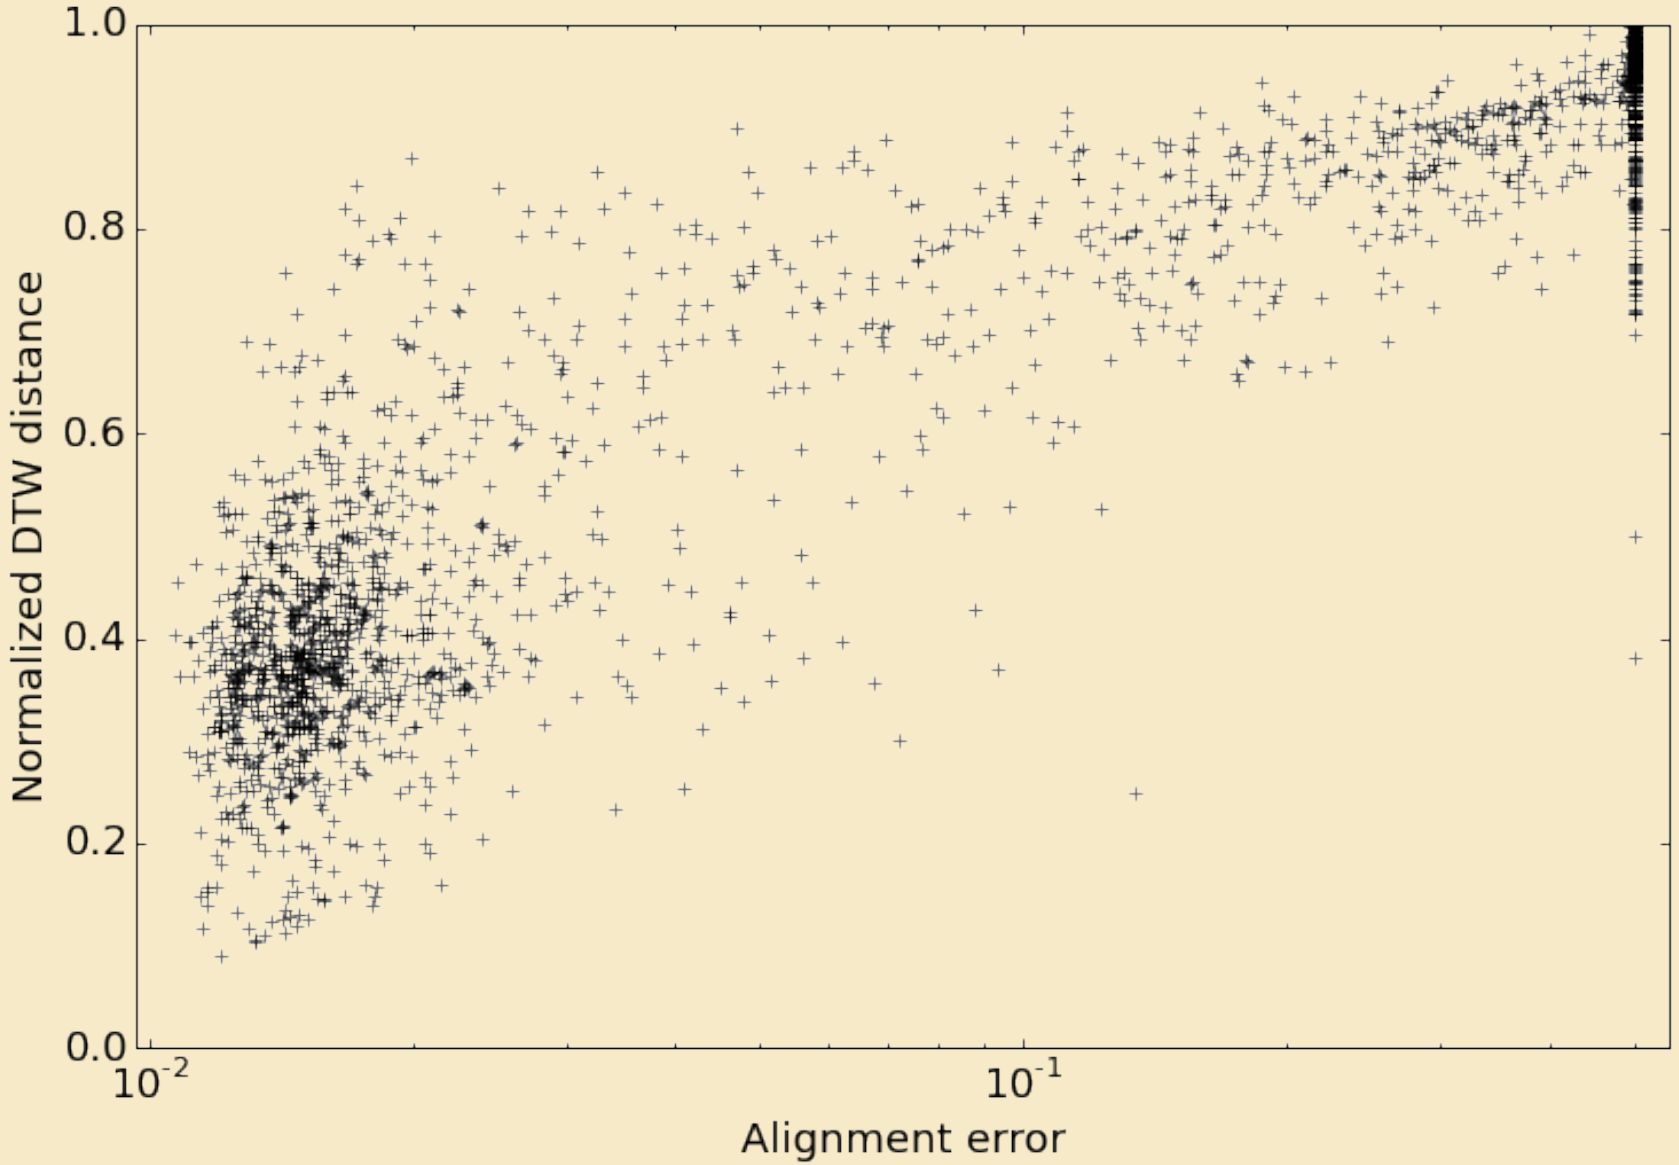

Best System:

- Use log-magnitude constant-Q spectrograms
- Don't beat synchronize
- L2 normalize spectra (cosine distance)
- Don't z-score spectrograms
- Use median distance as non-diagonal penalty
- Force sequences to match up to 96% of shorter
- Don't use a band path constraint
- Include penalties in confidence score
- Normalize by path length and submatrix mean

Real-World Test

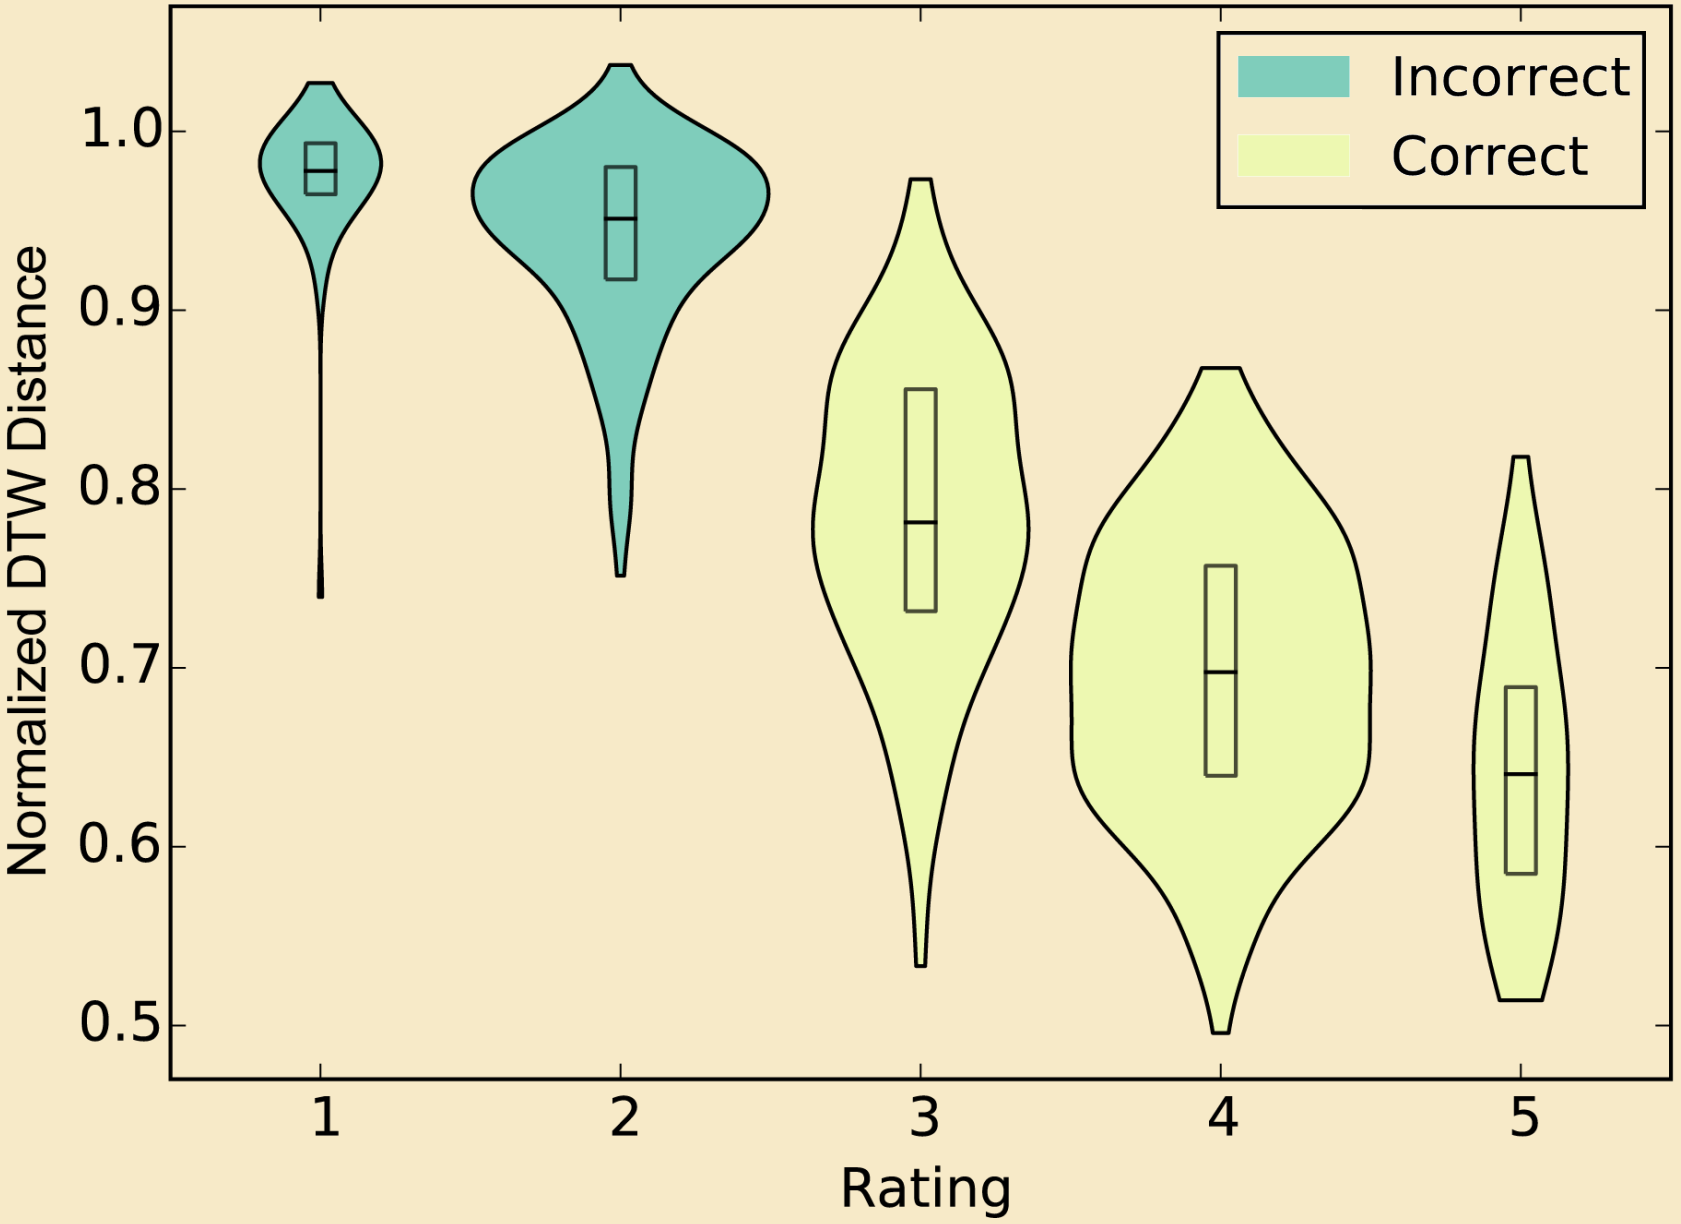

Pointers

http://bit.ly/alignment-overview
http://github.com/craffel/alignment-search
http://github.com/craffel/pretty-midi
http://github.com/craffel/djitw
http://github.com/bmcfee/librosa

craffel@gmail.com